

A Chain-of-Thought Prompting Approach with LLMs for Evaluating Students' Formative Assessment Responses in Science

Clayton Cohn¹, Nicole Hutchins¹, Tuan Le², Gautam Biswas¹

¹Vanderbilt University

²DePauw University

{clayton.a.cohn, nicole.m.hutchins, gautam.biswas}@vanderbilt.edu, tuanle_2025@depauw.edu

Abstract

This paper explores the use of large language models (LLMs) to score and explain short-answer assessments in K-12 science. While existing methods can score more structured math and computer science assessments, they often do not provide explanations for the scores. Our study focuses on employing GPT-4 for automated assessment in middle school Earth Science, combining few-shot and active learning with chain-of-thought reasoning. Using a human-in-the-loop approach, we successfully score and provide meaningful explanations for formative assessment responses. A systematic analysis of our method's pros and cons sheds light on the potential for human-in-the-loop techniques to enhance automated grading for open-ended science assessments.

Introduction

Improvements in Science, Technology, Engineering, and Mathematics (STEM) education have accelerated the shift from teaching and assessing facts to developing students' conceptual understanding and problem-solving skills (NGSS 2013). To foster students' developing scientific ideas and reasoning skills, it is crucial to have assessments that reveal and support their progress (Harris et al. 2023). Formative assessments play an important role in this endeavor, providing timely feedback and guidance when students face difficulties, which helps them to develop self-evaluation skills (Bloom, Madaus, and Hastings 1971). However, the process of grading and generating personalized feedback from frequent formative assessments is time-consuming for teachers and susceptible to errors (Rodrigues and Oliveira 2014; Haudek et al. 2011).

Large Language Models (LLMs) provide opportunities for automating short answer scoring (Funayama et al. 2023) and providing feedback to help students overcome their difficulties (Morris et al. 2023). These approaches can also aid teachers in identifying students' difficulties and generating actionable information to support student learning. To our knowledge, there is very little research that combines automated formative assessment grading and feedback generation for science domains where understanding, reasoning, and explaining are key to gaining a deep understanding of scientific phenomena (Mao et al. 2018).

Copyright © 2024, Association for the Advancement of Artificial Intelligence (www.aaai.org). All rights reserved.

This paper develops an approach for human-in-the-loop LLM prompt engineering using in-context learning and chain-of-thought reasoning with GPT-4 to support automated analysis and feedback generation for formative assessments in a middle school Earth Science curriculum. We present our approach, discuss our results, evaluate the limitations of our work, and then propose future research in this area of critical need in K-12 STEM instruction.

Background

To understand the difficulties students face when learning science, teachers need to actively track students' developing knowledge (Wiley et al. 2020). This is particularly important for open-ended, technology-enhanced learning environments that support students in their knowledge construction and problem-solving processes (Hutchins and Biswas 2023). In these environments, knowledge and skill development happen through system interactions that are difficult to monitor and interpret (Walkoe, Wilkerson, and Elby 2017). Formative assessments, evaluation, and feedback mechanisms aligned with target learning goals (Bloom, Madaus, and Hastings 1971), can play a dual role: (1) help students recognize constructs that are important to learning, and (2) provide teachers with a deeper understanding of student knowledge and reasoning to better support their developing STEM ideas (Cizek and Lim 2023). However, grading formative assessments, particularly in K-12 STEM contexts, where students' responses may not be well-structured and may vary considerably in vocabulary and stylistic expression, is time-consuming and can result in erroneous scoring and incomplete feedback (Liu et al. 2016). Moreover, grading these assessments at frequent intervals may become a burden rather than an aid to teachers. Very little research has examined effective mechanisms for generating automated grading and useful formative feedback for K-12 students that are aligned with classroom learning goals.

Advances in natural language processing (NLP) have produced improved automated assessment scoring approaches to support teaching and learning (e.g., Adair et al. 2023; Wilson et al. 2021). Proposed methodologies include data augmentation (Cochran, Cohn, and Hastings 2023), next sentence prediction (Wu et al. 2023), prototypical neural networks (Zeng et al. 2023), cross-prompt fine-tuning (Funayama et al. 2023), human-in-the-loop scoring via sam-

pling responses (Singla et al. 2022), and reinforcement learning (Liu et al. 2022). While these methods have enjoyed varying degrees of success, a majority of these applications have targeted more structured mathematics and computer science tasks (i.e., tasks that can be solved formulaically), but their grading is different from scoring free-form short-answer responses by middle school students in science domains. Data impoverishment concerns are common to educational data sets and a key consideration in applying these approaches to science assessments (Cochran et al. 2023). The data needed for training our models is small, imbalanced, and non-canonical in terms of syntax and semantics, all of which may impact model performance (Cohn 2020).

This research tackles several critical issues, namely: (1) grading open-ended, short-answer questions focused on science conceptual knowledge and reasoning, (2) utilizing LLMs to generate explanations aligned with specified learning objectives for both students and teachers and (3) addressing concerns related to data impoverishment. We hypothesize that our approach supports automated scoring and explanation that (1) aligns with learning objectives and standards, (2) provides actionable insight to students, especially in addressing their difficulties, and (3) engages teachers in the scoring and explanation generation process to resolve discrepancies and support the learning goals.

Methods

This section presents our curricular context, study design, dataset, LLM, and the details of our approach. Additional information regarding the formative assessment questions, rubrics, prompts, and method application can be found in the GitHub repository¹ along with test code and sample data.

Curricular context

This paper evaluates formative assessments conducted in the context of Science Projects Integrating Computing and Engineering (SPICE), an NGSS-aligned middle school earth sciences water runoff curriculum. Spanning three weeks, the curriculum tasks students with redesigning their schoolyard to enhance functionalities, using surface materials that minimize water runoff post-storm within specified cost and accessibility constraints (Chiu et al. 2019). We focus on formative assessments that are primarily linked to the conceptual understanding of water runoff and the conservation of matter principle (Hutchins et al. 2021).

Study Design and Dataset

This study utilized assessment data from two Vanderbilt University-approved SPICE studies involving 270 students at a Southeastern U.S. public middle school. Data was removed for non-consenting participants and some data was missing because of absences and incomplete submissions. We used evidence-centered design (ECD) (Mislevy and Haertel 2006) to align the assessments with the learning objectives of the SPICE curriculum.




Figure 1: The fictitious student’s conceptual model used by students to answers the assessment questions.

For this paper, we selected three questions that required students to analyze a pictorial model of water runoff (illustrated in Figure 1) and apply their conceptual knowledge and scientific reasoning to evaluate and explain the correct and incorrect components of the model. Each question was scored for at least one conceptual knowledge item, i.e., a correct application of a scientific fact. For example, in Q3, students had to identify that the arrow size representing total absorption was incorrect. Q2 and Q3 also required scoring students’ scientific reasoning, i.e., the use of scientific principles to explain an answer. For Q3, students could invoke the conservation principle to explain that the absorption arrow could not be larger than the rainfall arrow. The rubric assigned 1 point (conceptual) for Q1. Q2 and Q3 were scored for 4 points (2 items, 1 conceptual and 1 reasoning point for each item). For Q3, there were exactly 2 errors in the model. For Q2, students could choose from more than two correct phenomena, which resulted in differences in the grading results that we discuss later.

Model

Ever since OpenAI released ChatGPT² (a chatbot driven by the foundation model GPT-3.5) in November 2022, LLMs have received a tremendous amount of attention. Their ability to compose paper outlines, expository essays, and screenplays, has made the use of ChatGPT ubiquitous across academia, business, and news media. In March 2023, OpenAI released GPT-4³ (OpenAI 2023), which is largely considered the current state-of-the-art for LLMs (OpenAI 2023; Zhao et al. 2023). For this reason, we chose to use GPT-4 as the LLM to develop and evaluate our approach.

Approach

Brown et al. (2020) demonstrated that LLMs could “learn” from a few labeled instances in the prompt via *in-context learning* (ICL). Unlike fine-tuning, which requires expensive parameter updates and may result in decreased

¹<https://github.com/oele-isis-vanderbilt/EAAI24>

²<https://openai.com/blog/chatgpt>

³<https://openai.com/research/gpt-4>

performance for previously known tasks (Mosbach, Andriushchenko, and Klakow 2020), ICL uses the labeled instances in the prompt to generate text during inference that bypasses traditional training. This means that by simply changing the prompts, the same language model can be used across domains, tasks, and datasets without the need to modify the network’s parameters. Wei et al. (2022) extended this work by providing *chain-of-thought* (CoT) reasoning in the labeled instances. In contrast to a traditional ICL instance that only offers a question and its corresponding answer, CoT provides a reasoning chain with the answer. This helps the model generate correct inferences, and this reasoning is included in the model’s response along with the answer.

Eliciting reasoning is particularly useful for formative assessment scoring in science, where the open-ended nature of the questions can make scoring alignment difficult even between humans. Rather than generating a score only, CoT prompting elicits an explanation for the LLM’s response, enabling teachers to offer informed feedback to students. Alternatively, teachers can refine the rubric to improve grading for subsequent assessments. The model’s reasoning can also be used to identify specific causes of misalignment between the model and the teacher, which can then be leveraged to improve model output.

Active learning (Tan et al. 2023; Ren et al. 2021) takes a human-in-the-loop approach to improving model training, where the human as an “oracle” is consulted to label additional instances for inclusion in the next training iteration. By integrating CoT reasoning and active learning, educators or researchers can scrutinize instances with incorrect predictions to identify recurring patterns leading to the model’s errors across multiple instances. These patterns can be reintroduced into the prompt using CoT reasoning to rectify discrepancies between the model’s assessment and the human scorer. Moreover, combining CoT with active learning assists teachers and researchers in rectifying human errors in the initial scoring. This is particularly relevant when the humans confirm that the model’s scoring predictions are accurate.

We employ the inter-rater reliability (IRR) process to pinpoint scoring disagreements that may challenge the model, addressing them through CoT prompting. Active learning is then utilized to identify recurrent issues in the model’s alignment with the human scorers, and instances embodying these patterns are incorporated into the prompt with reasoning chains to correct the alignment. Once active learning concludes, the model is deployed for scoring new formative assessment responses through inference, accompanied by CoT reasoning to generate student feedback, and when needed, refining rubrics and formative assessment questions. Figure 2 provides a comprehensive overview of our approach.

Response Scoring. Two of this paper’s authors independently scored a randomly chosen 20% of the student responses for each of the three formative assessment questions using the rubric. Next, while conducting IRR, instances where the humans both agreed and disagreed on students’ scores were collected and included in the initial prompt.

Particular attention was paid to the misalignments between the graders that caused multiple instances to be scored differently before consensus was reached. To achieve consensus, the two reviewers discussed each scoring disagreement until they reached a consensus on how that particular instance should be scored. The agreed-upon instances acted as “ground truth” exemplars for the model to initially align itself with the human scorers. The instances where there were disagreements were used to pinpoint specific reasons for misalignment between the human scorers during IRR. We expected that the model might encounter the same misalignments during its scoring. This process was repeated for each of the three questions until Cohen’s $k > 0.7$ was achieved across all subscores for each question, after which one of this paper’s authors scored the full set of student responses. For this work, all students’ responses were manually graded to ensure accuracy while evaluating our method. Disagreements were resolved manually by the humans to form a consensus (described above). This consensus was used to align the LLM responses via CoT reasoning. In future work, as we collect more data, we will use the LLM to automatically score students’ responses and evaluate samples of the LLM’s generations to ensure accuracy. Furthermore, we refrained from updating the rubric during Active Learning; however, we intend to investigate this aspect in future research.

Before developing the initial prompt, we partitioned the dataset into training (80%) and testing (20%) instances for the three sets of formative assessment responses. The training set played a dual role in prompt development. Initially, few-shot examples were selected to construct the prompt, while the instances not utilized for few-shot learning served as a validation set for refining the prompt during active learning. Due to token limitations and the time cost for instance labeling, only a limited number of labeled instances were included in the prompt. As discussed in later sections, an excessive number of instances in the prompt can lead to overfitting. Furthermore, it is important that the validation set during active learning is sufficiently large to ensure accurate identification of scoring trends. In this paper, the validation-to-training set ratio was $\approx 43:1$.

Prompt Development. For prompting, we opted for the *persona pattern* (White et al. 2023), where the model was instructed to adopt the persona of a middle school teacher evaluating students’ formative assessment question responses. The prompt also provided the model with the formative assessment question and rubric, and the model was instructed to use the rubric to score students’ responses. The rubric also provided the model with the format to output its responses to improve readability and allow for programmatic parsing of the model’s generations.

Next, we incorporated ground truth examples into the prompt, complemented by CoT reasoning clarifying the reasons for awarding or not awarding points for each subscore. Following this, a comparable CoT input was included for instances where human scorers diverged in their assessments. This aimed at aligning the model with the IRR consensus, particularly when instances posed challenges similar to those faced by human reviewers in achieving consensus.

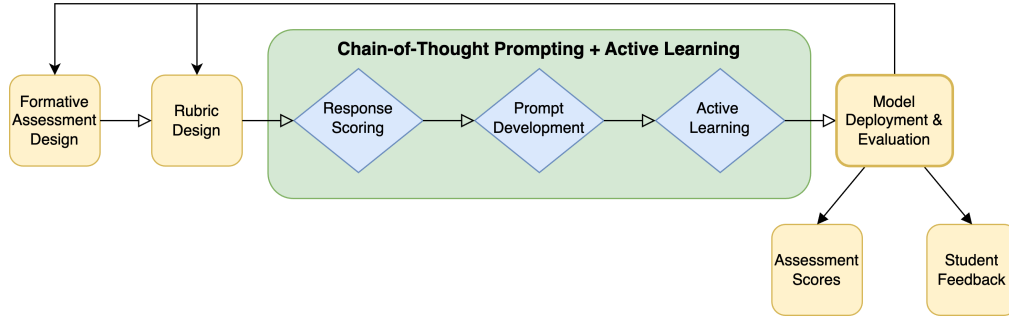


Figure 2: Our Chain-of-Thought Prompting + Active Learning approach. The green box encapsulates this process, where each of the blue diamonds is a step in that process. Yellow boxes represent the process’s application to the classroom.

For all labeled instances in the prompt, we used the following CoT reasoning template: evidence in the student’s response + reference to the rubric + score. We used quotations from the student’s response as evidence, tying it back to the rubric, and providing a score and explanation to the model; e.g., “The student says X. The rubric states Y. Based on the rubric, the student earned a score of Z.” This approach mirrored the original CoT publication (Wei et al. 2022), where algebraic word problems were broken down step-by-step to help the model arrive at the correct solution.

Additional labeled instances were added to the prompt as needed to balance the individual subscores. However, this was constrained by the small and imbalanced nature of our dataset. While investigating the effect that data balance has on the LLM’s performance is outside the scope of this work, in previous work (using a subset of the dataset used in this paper), we demonstrated that data balancing often improved language model performance (Cochran et al. 2022). For Q2 and Q3, balancing across 4 subscores was difficult, as adding one more instance to augment one subscore inherently affected the balance across the other subscores. Sometimes, achieving a perfect balance was not possible in the training set, but we included at least one positive and one negative instance across all subscores for each question’s prompt.

Active Learning. Validation set instances were fed through the model with the initial prompt and few-shot examples, and a researcher performed error analysis to discern patterns in the incorrect LLM generations. Specifically, we noted the reason for each incorrect scoring prediction and the faulty reasoning chains that caused the model to mislabel several instances. These reasoning chains were chosen as additional examples to add to the prompt, and CoT was used to correct the model’s reasoning errors. Candidate instances were prioritized for prompt inclusion based on the degree to which their reasoning errors caused other inaccurate model predictions, which resulted in correcting several wrongly scored instances.

There were only a few incorrectly predicted scores in the validation set for Q1, so all of those instances were added to the prompt during Active Learning. For Q2 and Q3, the researcher identified the n most useful instances, where n was defined as the minimum number of instances in the validation set that simultaneously addressed all of the LLM’s

reasoning errors and maintained data balance. This caused some overfitting, so we will experiment with 1-shot active learning to help mitigate this in future work. For all instances added to the prompt during Active Learning, we used CoT to correct the model’s faulty reasoning chains. We also re-balanced the few-shot instances across subscores during Active Learning to maintain data balance. In previous work, we showed that balancing training data to create a uniform label distribution can improve performance (Cochran et al. 2022). Other works have suggested balancing to achieve the true distribution of the dataset’s labels (Min et al. 2022).

In general, active learning can be performed until one of several stopping conditions is triggered: (1) the model achieves convergence, i.e., it no longer produces any incorrect validation scores; (2) the model predicts more validation scores incorrectly than in previous iterations, i.e., it overfits; and (3) there are not enough instances remaining in the validation set to achieve acceptable data balance in the prompt.

To test our method, we performed one iteration of active learning for each of the three formative assessment questions. For each subscore of a formative assessment question, we first identified scoring error trends, i.e., are model scoring errors mainly caused by false negatives (underscoring) or false positives (overscoring)? This alerted us to the “direction” in which we needed to guide the model to better align with the human scorers. We then examined the content of the incorrect validation set generations to identify common causes of incorrect scoring. We chose the most frequently occurring model reasoning error (i.e., the error that caused the model to wrongly predict the greatest number of validation set instances), and picked one of these instances to insert back into the prompt.

For example, with the *Runoff Arrow Direction* subscore in Q3, we found that the ratio of the model’s false positive to false negative predictions was 5:2. Additionally, we found that the cause of more than half of the false positives was due to the model awarding students a point for mentioning that the arrows in the diagram needed to change direction. This was incorrect because only the runoff arrow needed to change direction. To correct the model’s reasoning error, we chose one of the incorrect validation instances that included this reasoning error, inserted it into the prompt, and used CoT reasoning to help correct the model’s reasoning error

Q1 Arrow Size	n	Acc	F1	QWK
Zero-Shot	0	0.87	0.84	0.68
Few-Shot	4	1.00	1.00	1.00
Few-Shot, CoT	4	0.96	0.95	0.89
CoT + AL	12	0.98	0.97	0.95

Table 1: Performance comparisons for the Q1 *Arrow Size* subscore. For all questions, the best-performing scoring implementation is in bold for each metric, for each subscore (and total score). n refers to the number of few-shot instances used in the prompt.

going forward.

Results

We evaluated our method by comparing our model performance to the held-out test set across 4 implementations: three incremental baselines, and our Chain-of-Thought Prompting + Active Learning approach. We started with a *Zero-Shot* baseline, where the rubric is included in the prompt, but no labeled examples were present. We then used a *Few-Shot* baseline, where we provided the model with labeled instances in the prompt, but the labeled instances only consisted of numerical scores (i.e., no CoT reasoning). Our third and final baseline, *Few-Shot, CoT*, added CoT reasoning to the few-shot instances. Last, we employed our Chain-of-Thought Prompting + Active Learning approach and compared it to the three baselines. Evaluating our approach across these incremental baselines allowed us to examine the effects of adding specific parts of the pipeline and to understand the degree to which each component contributed to the model’s ability to score and explain formative assessment responses.

To compare implementations, we chose the Macro F1-Score and Cohen’s Quadratic Weighted Kappa (QWK) (Cohen 1968) metrics. The F1-Score is prevalent in the literature for evaluating overall model performance. Macro F1 was chosen, specifically, due to our dataset’s imbalance across subscores. Often, scientific reasoning subscores are heavily weighted towards the negative class (i.e., a large majority of the students do not demonstrate scientific reasoning). Cohen’s QWK was chosen because it is widely used in the automated essay scoring (AES) literature (Singh et al. 2023; Singla et al. 2022). Unlike traditional Cohen’s k (Cohen 1960), Cohen’s QWK accounts for the degree of disagreement, making it well-suited for ordinal data. We included accuracy for reference, but we do not use it in our actual performance comparisons.

Model performance comparisons for each of the three formative assessment questions are shown in Tables 1, 2, and 3.

Question 1: Q1 asked students what the different-sized arrows in the diagram meant. A student received a point for correctly identifying that the diagram used the size of the arrows to represent the quantity of water (concept: “Arrow Size”, see Table 1).

Q1 took 2 rounds of IRR for the human scorers to reach a consensus. The grading involved scoring for one possible

Q2 Arrow Direction	n	Acc	F1	QWK
Zero-Shot	0	0.91	0.89	0.78
Few-Shot	5	0.87	0.79	0.60
Few-Shot, CoT	5	0.98	0.98	0.95
CoT + AL	10	0.98	0.98	0.95
Q2 Arr. Dir., Reasoning	n	Acc	F1	QWK
Zero-Shot	0	0.92	0.73	0.47
Few-Shot	5	0.89	0.67	0.36
Few-Shot, CoT	5	0.91	0.70	0.41
CoT + AL	10	0.92	0.65	0.3
Q2 Arrow Size	n	Acc	F1	QWK
Zero-Shot	0	0.77	0.69	0.39
Few-Shot	5	0.77	0.69	0.39
Few-Shot, CoT	5	0.91	0.88	0.77
CoT + AL	10	0.94	0.92	0.83
Q2 Arr. Sz., Reasoning	n	Acc	F1	QWK
Zero-Shot	0	0.96	0.82	0.65
Few-Shot	5	0.98	0.90	0.79
Few-Shot, CoT	5	0.94	0.77	0.55
CoT + AL	10	0.96	0.82	0.65
Q2 Total Score	n	Acc	F1	QWK
Zero-Shot	0	0.60	0.59	0.65
Few-Shot	5	0.53	0.52	0.55
Few-Shot, CoT	5	0.75	0.80	0.80
CoT + AL	10	0.85	0.79	0.87

Table 2: Performance comparisons for Question 2.

point and no science reasoning subscores. GPT-4 aligned with the human scorer to a “moderate” degree (QWK ≥ 0.6) (McHugh 2012) even in a zero-shot setting. Once labeled instances were added, the model achieved a perfect score on the test set. When CoT reasoning was provided, performance decreased for both Macro F1 and QWK, as the reasoning chains initially caused the model to deviate from the human scorer. Once active learning was performed, however, much of that performance gap was closed due to the additional few-shot instances and model reasoning error corrections.

Question 2: Q2 asked students to identify two things that the diagram did well for 4 possible points: 2 for science concepts, and 2 for science reasoning. For the science concepts subscores, the student received a point for *Arrow Direction* if he or she correctly identified that the diagram did a good job of showing that water originated from the sky in the form of rain, some water was absorbed, or some resulted in runoff. For *Arrow Size*, students received a point if they discussed that the diagram did a good job of using arrow size to represent the water amount. Each of the *Arrow Direction* and *Arrow Size* subscores also included an additional possible point if students demonstrated scientific reasoning in their responses (see Table 2).

Q2 science concepts subscores (i.e., *Arrow Direction* and *Arrow Size*) saw their best performance (or tied for best performance) using the full Chain-of-Thought Prompting + Active Learning method. The science reasoning subscores’

performances decreased as additional components of the method were added, i.e., after CoT and active learning were introduced. Overall, the total score was best when the complete method was used, as this resulted in the highest QWK value.

Q2 was the most difficult for the human scorers to agree on and it required three separate IRR rounds to achieve consensus. Some of the difficulty in scoring may be attributed to the open-ended nature of the question. There are multiple ideas in the conceptual model that are correct, but students were only asked to identify two things the diagram explained well. Many students responded vaguely, and several students provided both correct and incorrect statements in the same response. These types of ambiguous instances were difficult for the human scorers to agree on even when they awarded points based on the rubric. It seems the LLM encountered the same types of issues.

Consider a student whose Q2 response was, “[the arrow represents] the amount of absorption”. Arguably, the student understood that the model’s arrows corresponded to the quantity of water. However, the absorption arrow in the diagram was incorrect (it was larger than the rainfall arrow, so it violated the law of conservation of matter). Because the question asked for examples of things the diagram does a *good* job of, and the absorption arrow was incorrect, both reviewers felt responses like this one should not receive a point for *Arrow Size* even though the student may understand that arrow size corresponds to the amount of water. During our active learning validation run, the model incorrectly awarded several points to these types of responses. When we attempted to use CoT to correct the model’s reasoning error, the model began to mislabel other instances it had previously scored correctly, i.e., there was overfitting. Ultimately, the researchers agreed that both the Q2 question wording and the Q2 rubric need to be rewritten to provide clearer guidance to the students.

Question 3: Q3 asked students to list two erroneous things they would change in the conceptual model diagram. Like Q2, 4 total points were assigned to Q3: 2 for science concepts and 2 for scientific reasoning. The science concepts subscores were: (1) *Runoff Direction*, where the student received a point if he or she indicated that the runoff arrow was pointing the wrong direction (uphill) and (2) *Arrow Size*, where a point was awarded if the student mentioned that the arrow sizes needed to change and adhere to the law of conservation of matter. Similar to Q2, students got additional points if they demonstrated correct scientific reasoning in their responses (see Table 3).

All Q3 subscores (science concepts and scientific reasoning) improved performance across both metrics (except Macro F1 for total score) after we added the few-shot examples. When CoT was added, performance increased for both *Runoff Direction* subscores but decreased substantially for both *Arrow Size* subscores. We saw similar behavior in the Q1 *Arrow Size* subscore, where adding CoT reasoning caused the model to become misaligned with the human. Once the Active Learning component was added, however, all subscores except *Runoff Direction* achieved

Q3 Runoff Direction	n	Acc	F1	QWK
Zero-Shot	0	0.89	0.88	0.77
Few-Shot	5	0.91	0.90	0.80
Few-Shot, CoT	5	0.92	0.92	0.84
CoT + AL	9	0.89	0.88	0.75
Q3 Run. Dir., Reasoning	n	Acc	F1	QWK
Zero-Shot	0	0.94	0.89	0.79
Few-Shot	5	0.94	0.91	0.82
Few-Shot, CoT	5	0.94	0.92	0.83
CoT + AL	9	0.98	0.97	0.94
Q3 Arrow Size	n	Acc	F1	QWK
Zero-Shot	0	0.87	0.83	0.67
Few-Shot	5	0.89	0.87	0.73
Few-Shot, CoT	5	0.85	0.83	0.65
CoT + AL	9	0.92	0.92	0.83
Q3 Arr. Sz., Reasoning	n	Acc	F1	QWK
Zero-Shot	0	0.98	0.90	0.79
Few-Shot	5	1.00	1.00	1.00
Few-Shot, CoT	5	0.94	0.82	0.64
CoT + AL	9	1.00	1.00	1.00
Q3 Total Score	n	Acc	F1	QWK
Zero-Shot	0	0.74	0.80	0.85
Few-Shot	5	0.75	0.73	0.87
Few-Shot, CoT	5	0.75	0.71	0.79
CoT + AL	9	0.81	0.80	0.90

Table 3: Performance comparisons for Question 3.

their best performance across both metrics. *Runoff Direction* achieved its best performance when CoT was added, but was overfit during active learning. Unlike Q2, where the best-performing subscores were the science concepts, both science reasoning subscores did better than their science concepts counterparts for both metrics.

For Q3, the human scorers achieved consensus quickly after 1 round of IRR, and only one issue caused multiple scoring disagreements. The model’s reasoning errors for the scientific reasoning subscores were easily addressed via CoT (relative to Q2). A major model reasoning error for Q3 was that it tended to cite the same piece of evidence to justify awarding points for different subscores (i.e., it overscored). This was a disagreement with the human scorers, but we did not include it in the initial prompt or few-shot CoT reasoning chains. Once this model reasoning error was addressed during active learning, the issue was largely mitigated, and performance improved across the board.

Summary: Across all questions, the model’s scoring mostly aligned with the human scorers. Of the 11 subscores and total scores, 9 of them saw “strong” agreement or better (QWK ≥ 0.8) at some point in the process (i.e., across the 4 implementations: 3 baselines and our Chain-of-Thought Prompting + Active Learning approach). 4 subscores achieved “almost perfect” (QWK > 0.9) agreement. All subscores except one (Q2 Arrow Direction Reasoning) saw a Macro F1 of 0.90 or greater at some point in the process.

Importantly, we also demonstrated that both CoT reasoning and active learning run the risk of overfitting, particularly when applied to the less complex science concepts questions (e.g., Q1 *Arrow Size* and Q3 *Runoff Direction*) and the more ambiguous scientific reasoning questions (e.g., Q2 *Arrow Direction Reasoning* and Q2 *Arrow Size Reasoning*). It should also be noted that the level of agreement during IRR may provide a ballpark expectation of model performance, as we found questions that were easier for the human scorers to agree on were also easier for the model to correctly align with the human scorers. Similarly, in questions where the human scorers had difficulty achieving consensus, the model had difficulty with scoring. More research needs to be done to evaluate this quantitatively.

Comparing Model and Human Performance

We applied inductive coding (Charmaz 2006) to evaluate performance and identify future directions to improve our human-in-the-loop approach. First, the lead author (not involved in rubric creation and scoring) reviewed all instances in which the model and the human coder disagreed and identified agreement with the model in 3 out of the 22 disagreements (1 conceptual disagreement, 2 reasoning disagreements). The research team reviewed the results to evaluate what may have caused scoring errors and to identify potential future directions for improvement. During the review process, the team created memos of key findings (Hatch 2002). The team compared the memos and came up with three key themes for improvement in future work:

1. **Need for Additional Mechanisms to Target Model Deficiencies:** Differences in scoring identified that the model showed a tendency to overfit in some cases. For instance, if the CoT got too granular, the model demonstrated issues that were related to keywords such as “because” (e.g., the model identified it as a demonstration of reasoning), “arrow size” (e.g., the model assumed that use of the terminology indicated a correct application even if correct attributions were not made to the scientific process), and vocabulary definitions (e.g., the model did not realize “run off” and “runoff” were identical). In a small set of cases, the model cited a student’s faulty logic to justify awarding a point for a response and reused the same piece of evidence to award points for both concept identification and reasoning;
2. **Ability to Leverage the Model to Support Rubric Refinements:** Comparing reviewer and model differences for Q2 helped identify limitations in the original rubric for such an open-ended question. Utilizing the results and the explanations provided by the model, this human-in-the-loop approach can benefit teachers and researchers in refining the rubrics and scoring mechanisms to better support instruction and student learning; and
3. **Resolve Unexplained Model Applications:** In some cases, the model did not follow CoT reasoning and did not provide evidence of its positive predictions even though all positive prompt instances provided this evidence. This may be a potential limitation in the approach to providing feedback for positive performances.

Overall, our approach was successful, but the instances discussed above provide opportunities for future work to improve model output, rubric development, and sometimes even reworking questions to make them clearer.

Conclusion and Future Implications

In this paper, we employed a Chain-of-Thought Prompting + Active Learning approach for scoring and explaining formative assessment question responses in a middle school Earth Science curriculum. Our results show that GPT-4, CoT reasoning, and active learning can be effectively leveraged toward accurate grading of science formative assessments. In several cases, the model achieved “almost perfect” alignment with humans. The model generated relevant evidence linked to the rubric to help explain its scoring, which could benefit students and teachers. We also analyzed the model’s weaknesses and identified several areas for improving LLM-based assessments.

Limitations: With LLM approaches, ethical concerns arise with regard to privacy, bias, and hallucinations (Zhuo et al. 2023), and these concerns are amplified when they are deployed in high-stakes environments (e.g., classrooms with children). In addition, while CoT has been shown to improve model performance over traditional ICL, the degree to which the reasoning chains guide the model’s decision-making (if at all) is still an open question (Turpin et al. 2023). Our results also show that CoT and active learning can lead to overfitting, in particular, with simpler, easier-to-define sub-problems. In these cases, LLM approaches may be overkill, as Moore et al. (2023) recently demonstrated. Rule-based methods outperformed GPT-4 in detecting common item-writing flaws in student-generated multiple-choice questions.

Looking to the Future: Anecdotally, in an interview with middle school science teachers who implemented the curriculum, the teachers identified the potential benefits of these explanations as tools to inform students on where to go next in their learning, as opposed to assigning performance scores. We aim to extend this partnership with classroom teachers to mold the LLM’s output to best fit their needs, and investigate how we can best use our method to evaluate students’ learning performance and improve students’ learning. As we continue to refine our approach, we hope these enhancements will pave the way for more effective and efficient LLM applications in science education.

Acknowledgments

This work is supported by the National Science Foundation under awards DRL-2112635 and IIS-2017000. Any opinions, findings, conclusions, and recommendations in this paper are those of the authors and do not necessarily reflect the views of the National Science Foundation.

References

Adair, A.; Pedro, M. S.; Gobert, J.; and Segan, E. 2023. Real-Time AI-Driven Assessment and Scaffolding that Improves Students’ Mathematical Modeling during Science Investigations. In Wang, N.; Rebollo-Mendez, G.; Mat-

- suda, N.; Santos, O. C.; and Dimitrova, V., eds., *Artificial Intelligence in Education*, 202–216. Cham: Springer Nature Switzerland. ISBN 978-3-031-36272-9.
- Bloom, B.; Madaus, G.; and Hastings, J. 1971. *Handbook on Formative and Summative Evaluation of Student Learning*. New York: McGraw-Hill.
- Brown, T. B.; Mann, B.; Ryder, N.; Subbiah, M.; Kaplan, J.; Dhariwal, P.; Neelakantan, A.; Shyam, P.; Sastry, G.; Askell, A.; Agarwal, S.; Herbert-Voss, A.; Krueger, G.; Henighan, T.; Child, R.; Ramesh, A.; Ziegler, D. M.; Wu, J.; Winter, C.; Hesse, C.; Chen, M.; Sigler, E.; Litwin, M.; Gray, S.; Chess, B.; Clark, J.; Berner, C.; McCandlish, S.; Radford, A.; Sutskever, I.; and Amodei, D. 2020. Language Models are Few-Shot Learners. *arXiv e-prints*, arXiv:2005.14165.
- Charmaz, K. 2006. *Constructing grounded theory: A practical guide through qualitative analysis*. Sage.
- Chiu, J.; McElhaney, K.; Zhang, N.; Biswas, G.; Fried, R.; Basu, S.; Alozie, N.; and Hong, J. 2019. A Principled Approach to NGSS-aligned Curriculum Development Integrating Science, Engineering, and Computation: A Pilot Study. In *NARST Annual International Conference*. NARST.
- Cizek, G. J.; and Lim, S. N. 2023. Formative assessment: an overview of history, theory and application. In Tierney, R. J.; Rizvi, F.; and Ercikan, K., eds., *International Encyclopedia of Education (Fourth Edition)*, 1–9. Oxford: Elsevier, fourth edition edition. ISBN 978-0-12-818629-9.
- Cochran, K.; Cohn, C.; and Hastings, P. 2023. Improving NLP model performance on small educational data sets using self-augmentation. In *Proceedings of the 15th International Conference on Computer Supported Education (2023, to appear)*.
- Cochran, K.; Cohn, C.; Hutchins, N.; Biswas, G.; and Hastings, P. 2022. Improving automated evaluation of formative assessments with text data augmentation. In *International Conference on Artificial Intelligence in Education*, 390–401. Springer.
- Cochran, K.; Cohn, C.; Rouet, J. F.; and Hastings, P. 2023. Improving Automated Evaluation of Student Text Responses Using GPT-3.5 for Text Data Augmentation. In *International Conference on Artificial Intelligence in Education*, 217–228. Springer.
- Cohen, J. 1960. A coefficient of agreement for nominal scales. *Educational and psychological measurement*, 20(1): 37–46.
- Cohen, J. 1968. Weighted kappa: nominal scale agreement provision for scaled disagreement or partial credit. *Psychological bulletin*, 70(4): 213.
- Cohn, C. 2020. *BERT efficacy on scientific and medical datasets: a systematic literature review*. DePaul University.
- Funayama, H.; Asazuma, Y.; Matsubayashi, Y.; Mizumoto, T.; and Inui, K. 2023. Reducing the Cost: Cross-Prompt Pre-finetuning for Short Answer Scoring. In *International Conference on Artificial Intelligence in Education*, 78–89. Springer.
- Harris, C.; Wiebe, E.; Grover, S.; and Pellegrino, J. 2023. *Classroom-based STEM assessment: Contemporary issues and perspectives*. Community for Advancing Discovery Research in Education (CADRE). Education Development Center, Inc.
- Hatch, J. A. 2002. *Doing qualitative research in education settings*. SUNY Press.
- Haudek, K. C.; Kaplan, J. J.; Knight, J.; Long, T.; Merrill, J.; Munn, A.; Nehm, R.; Smith, M.; and Urban-Lurain, M. 2011. Harnessing technology to improve formative assessment of student conceptions in STEM: forging a national network. *CBE—Life Sciences Education*, 10(2): 149–155.
- Hutchins, N.; and Biswas, G. 2023. Using Teacher Dashboards to Customize Lesson Plans for a Problem-Based, Middle School STEM Curriculum. In *LAK23: 13th International Learning Analytics and Knowledge Conference*, LAK2023, 324–332. New York, NY, USA: Association for Computing Machinery. ISBN 9781450398657.
- Hutchins, N. M.; Basu, S.; McElhaney, K.; Chiu, J.; Fick, S.; Zhang, N.; and Biswas, G. 2021. Coherence across conceptual and computational representations of students’ scientific models. In *The International Society of the Learning Sciences Annual Meeting 2021*. International Society of the Learning Sciences (ISLS).
- Liu, E.; Stephan, M.; Nie, A.; Piech, C.; Brunskill, E.; and Finn, C. 2022. Giving Feedback on Interactive Student Programs with Meta-Exploration. *Advances in Neural Information Processing Systems*, 35: 36282–36294.
- Liu, O. L.; Rios, J. A.; Heilman, M.; Gerard, L.; and Linn, M. C. 2016. Validation of automated scoring of science assessments. *Journal of Research in Science Teaching*, 53(2): 215–233.
- Mao, L.; Liu, O. L.; Roohr, K.; Belur, V.; Mulholland, M.; Lee, H.-S.; and Pallant, A. 2018. Validation of automated scoring for a formative assessment that employs scientific argumentation. *Educational Assessment*, 23(2): 121–138.
- McHugh, M. L. 2012. Interrater reliability: the kappa statistic. *Biochemia medica*, 22(3): 276–282.
- Min, S.; Lyu, X.; Holtzman, A.; Artetxe, M.; Lewis, M.; Hajishirzi, H.; and Zettlemoyer, L. 2022. Rethinking the role of demonstrations: What makes in-context learning work? *arXiv preprint arXiv:2202.12837*.
- Mislevy, R. J.; and Haertel, G. D. 2006. Implications of Evidence-Centered Design for Educational Testing. *Educational Measurement: Issues and Practice*, 25(4): 6–20.
- Moore, S.; Nguyen, H. A.; Chen, T.; and Stamper, J. 2023. Assessing the Quality of Multiple-Choice Questions Using GPT-4 and Rule-Based Methods. In *European Conference on Technology Enhanced Learning*, 229–245. Springer.
- Morris, W.; Crossley, S.; Holmes, L.; Ou, C.; McNamara, D.; and Dascalu, M. 2023. Using Large Language Models to Provide Formative Feedback in Intelligent Textbooks. In *International Conference on Artificial Intelligence in Education*, 484–489. Springer.
- Mosbach, M.; Andriushchenko, M.; and Klakow, D. 2020. On the stability of fine-tuning bert: Misconceptions, explanations, and strong baselines. *arXiv preprint arXiv:2006.04884*.

- NGSS. 2013. *Next Generation Science Standards: For States, By States*. The National Academies Press.
- OpenAI. 2023. GPT-4 Technical Report. *arXiv e-prints*, arXiv:2303.08774.
- Ren, P.; Xiao, Y.; Chang, X.; Huang, P.-Y.; Li, Z.; Gupta, B. B.; Chen, X.; and Wang, X. 2021. A survey of deep active learning. *ACM computing surveys (CSUR)*, 54(9): 1–40.
- Rodrigues, F.; and Oliveira, P. 2014. A system for formative assessment and monitoring of students’ progress. *Computers & Education*, 76: 30–41.
- Singh, S.; Pupneja, A.; Mital, S.; Shah, C.; Bawkar, M.; Gupta, L. P.; Kumar, A.; Kumar, Y.; Gupta, R.; and Shah, R. R. 2023. H-AES: Towards Automated Essay Scoring for Hindi. *arXiv preprint arXiv:2302.14635*.
- Singla, Y. K.; Krishna, S.; Shah, R. R.; and Chen, C. 2022. Using sampling to estimate and improve performance of automated scoring systems with guarantees. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 36 (11), 12835–12843.
- Tan, W.; Lin, J.; Lang, D.; Chen, G.; Gašević, D.; Du, L.; and Buntine, W. 2023. Does informativeness matter? Active learning for educational dialogue act classification. In *International Conference on Artificial Intelligence in Education*, 176–188. Springer.
- Turpin, M.; Michael, J.; Perez, E.; and Bowman, S. R. 2023. Language Models Don’t Always Say What They Think: Unfaithful Explanations in Chain-of-Thought Prompting. *arXiv preprint arXiv:2305.04388*.
- Walkoe, J.; Wilkerson, M.; and Elby, A. 2017. Technology-Mediated Teacher Noticing: A Goal for Classroom Practice, Tool Design, and Professional Development. In *Proceedings of the 12th International Conference on Computer Supported Collaborative Learning (CSCL) 2017*. International Society of the Learning Sciences.
- Wei, J.; Wang, X.; Schuurmans, D.; Bosma, M.; Ichter, B.; Xia, F.; Chi, E.; Le, Q.; and Zhou, D. 2022. Chain-of-Thought Prompting Elicits Reasoning in Large Language Models. *arXiv e-prints*, arXiv:2201.11903.
- White, J.; Fu, Q.; Hays, S.; Sandborn, M.; Olea, C.; Gilbert, H.; Elnashar, A.; Spencer-Smith, J.; and Schmidt, D. C. 2023. A prompt pattern catalog to enhance prompt engineering with chatgpt. *arXiv preprint arXiv:2302.11382*.
- Wiley, K. J.; Dimitriadis, Y.; Bradford, A.; and Linn, M. C. 2020. From Theory to Action: Developing and Evaluating Learning Analytics for Learning Design. In *Proceedings of the Tenth International Conference on Learning Analytics & Knowledge, LAK ’20*, 569–578. New York, NY, USA: Association for Computing Machinery. ISBN 9781450377126.
- Wilson, J.; Ahrendt, C.; Fudge, E. A.; Raiche, A.; Beard, G.; and MacArthur, C. 2021. Elementary teachers’ perceptions of automated feedback and automated scoring: Transforming the teaching and learning of writing using automated writing evaluation. *Computers & Education*, 168: 104208.
- Wu, X.; He, X.; Liu, T.; Liu, N.; and Zhai, X. 2023. Matching exemplar as next sentence prediction (mensp): Zero-shot prompt learning for automatic scoring in science education. In *International Conference on Artificial Intelligence in Education*, 401–413. Springer.
- Zeng, Z.; Li, L.; Guan, Q.; Gašević, D.; and Chen, G. 2023. Generalizable Automatic Short Answer Scoring via Prototypical Neural Network. In *International Conference on Artificial Intelligence in Education*, 438–449. Springer.
- Zhao, W. X.; Zhou, K.; Li, J.; Tang, T.; Wang, X.; Hou, Y.; Min, Y.; Zhang, B.; Zhang, J.; Dong, Z.; et al. 2023. A survey of large language models. *arXiv preprint arXiv:2303.18223*.
- Zhuo, T. Y.; Huang, Y.; Chen, C.; and Xing, Z. 2023. Red teaming chatgpt via jailbreaking: Bias, robustness, reliability and toxicity. *arXiv preprint arXiv:2301.12867*, 12–2.